

Recommender System

Hugo Veríssimo

Foundations of Machine Learning 24/25
University of Aveiro
Aveiro, Portugal
hugoverissimo@ua.pt

João Cardoso

Foundations of Machine Learning 24/25
University of Aveiro
Aveiro, Portugal
joaopcardoso@ua.pt

Abstract—

Keywords: MovieLens, GroupLens, Recommender System, Collaborative Filtering

I. INTRODUCTION

Since the early 90's that the production rate of multimedia content has increased dramatically (pun intended). Initially, users relied mostly on video store owners, film critics on newspapers, friends. With increasing volumes of content, recommender systems have appeared as data-driven methods to reliably and quickly recommend movies based on the users' own appreciations.

Recommender systems are information tools that provide users with recommended items (ideally) based on their list of preferences [1], [2]. These can be divided in different types, depending on the algorithm and approach to the data. Three categories can be considered: non-personalized, content-based, and collaborative filtering algorithms. Regardless of the algorithm of choice, these face crucial challenges: cold start problem (where it is difficult to tailor recommendations to a user without known preferences, or recommend an item with no reviews); data sparsity (given that most users review few items in the universe of possible items, leaving most of the user/item matrix empty); and scalability (as data grows exponentially, processing becomes evermore expensive and troublesome). These systems have been widely used in many different areas (online shopping, music, books, movie recommendation), and significant investment has gone into developing evermore personalized algorithms. A notable case for this was the Netflix prize competition in 2006, a moment where the research in the field skyrocketed. As a result, it was needed to develop better frameworks for comparison, considering not only the metrics, but data preprocessing, preparation, and routine, in order to ensure reproducibility across models and authors [3].

For that reason, the developed recommender systems is based on one of the most widely used movie databases, MovieLens dataset, for education and development [4]. The choice of this dataset allowed for a based comparison with algorithms from the literature, and facilitate the analysis and interpretation of the results here presented.

II. STATE OF THE ART

The field of recommender systems is wide and covers many different algorithms and techniques. In this paper we focus on collaborative filtering using matrix factorization, as it is the main focus of the work here developed.

The Netflix prize is often cited as one of the main drivers for research in collaborative filtering recommender systems [5]. One of the initial awarded proposals was that by Brandyn Webb, known by his alias "Simon Funk". Despite the name of the algorithm (FunkSVD), Singular Value Decomposition is not used, it uses instead gradient descent to find the latent feature values used to predict the ratings matrix. The algorithm uses only the available ratings, representing a great advantage against certain SVD methods.

$$r_{ij} = u_i \cdot v_j \quad (1)$$

$$u_{if}(\text{new}) = u_{if}(\text{old}) + 2\alpha(r_{ij} - \tilde{r}_{ij})v_{jf} \quad (2)$$

$$v_{jf}(\text{new}) = v_{jf}(\text{old}) + 2\alpha(r_{ij} - \tilde{r}_{ij})u_{if} \quad (3)$$

There is a caveat however, given that the model requires fine tuning of numerous parameters (number of latent features, learning rate, training iterations, regularization parameter), which can lead to overfitting [6]. In the work by Zhou et al., alternating least squares with weighted λ regularization (ALS-WR).

$$f(U, M) = \sum_{\{i,j\} | r_{i,j} \in I} (r_{i,j} - u_i^T m_j)^2 + \lambda \left(\sum_i n_{u_i} \|u_i\|^2 + \sum_j n_{m_j} \|m_j\|^2 \right) \quad (4)$$

The algorithm expresses the rating matrix as the product of two smaller matrices U (user matrix) and M (item matrix). Thanks to its simplicity, the algorithm tackles both scalability and sparseness of user profiles, with the added bonus of not overfitting [7]. A third model by Gopalan et al. uses a probabilistic approach by assuming that the observed rating is drawn from a Poisson distribution, which is parameterized by the inner product of a user weights vector and an item weights vector.

$$y_{ui} \sim \text{Poisson}(\theta_u^T \beta_i) \quad (5)$$

Where:

User weights: $\theta_u = [\theta_{u1}, \dots, \theta_{uk}]$

Item weights: $\beta_i = [\beta_{i1}, \dots, \beta_{ik}]$

$\theta_{uk} \sim \text{Gamma}(a, b)$

$\beta_{ik} \sim \text{Gamma}(c, d)$

With this, the model is able to compute how likely the user is to consume new items [8].

III. METHODOLOGY

A. Data description

The dataset MovieLens for Education and Research (small) was used to test the different models. It contains 100 836 ratings, from 1 to 5, for 9724 movies (and its genres) and 610 users, where each user rated at least 20 movies [9].

B. Data splitting & models implemented

The dataset was split between train and test (80/20), which resulted in 80668 ratings for the training and 20168 for the testing dataset. Three models were developed, as described in the Table I.

TABLE I: Models implemented to the MovieLens dataset.

Models - features	Description
Model 01 - Users and Movies	Collaborative filtering, Linear Regression (CF-LR)
Model 02 - Users and Movies + Genres	Collaborative filtering, Linear Regression content based (CF-LRC)
Model 03 - Users and Movies + Genres	Collaborative filtering, Single Value Decomposition (CF-SVD)

Our models were fitted with an 8-fold cross validation to find the best hyperparameters for the considered interval, as described in the lectures. For Model 03, we used the package *surprise* and the tools therein [10].

The error metrics used in the present work were the mean absolute error (MAE) and the root mean squared error (RMSE). The MAE considers all error equally, regardless of their size, whereas RMSE strongly penalizes larger errors.

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| \quad (6)$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (7)$$

C. Exploratory data analysis

The dataset consists of several movies across a wide range of genres. As seen in Fig. 1 there is positively skewed distribution, as is common in this type of dataset.




Fig. 1: Distribution of movies per movie genre.

Despite the variety in genres, it's important to retain that generally movies are more complex than representing a single genre. In that sense, it is useful to assess the similarity between genres (using the cosine similarity index), as illustrated in Fig. 2. This helps understand how the task of recommending a movie can start to grow more complex as we add more and more detail to the dataset. In Table II the five most related genres are shown.




Fig. 2: Cosine similarity index between genres.

TABLE II: The five most related genres per cosine similarity index.

Genre	Genre	Similarity index
Animation	Children	.47
Action	Adventure	.40
Crime	Thriller	.38
Romance	Comedy	.36
Romance	Drama	.35

The dataset is relatively recent, with movies from the 2010's, going all the way back to the early 1900's, as per Fig. 3. As expected, there are much less ratings for older movies, for two main reasons: older movies are less popular and there are less movies in general.




Fig. 3: Distribution of movies per year.

As expected, both distributions from Fig. 4 and 5 are positively skewed as mentioned before. In general there are more people rating a small amount of movies, and the number of movies with higher counts of ratings tends to decrease rapidly (as there are very few, very popular movies).

This type of distribution is commonly found across recommender systems, and represents well the challenges posed to these algorithms: sparsity, whereas most of the matrix is empty (98.3 % of missing values); bias, popular movies tend to dominate the analysis; and the long tail represents the challenge to make recommendations based on the large number of movies (or users) with very few ratings (or that rated few movies).




Fig. 4: Distribution of the number of movies and the amount of ratings. The first bin correspond to 1 movie rated.

To finalize, it is important to mention (again) that the minimum number of ratings per user is 20, and the maximum is 2698. The minimum rating per movie is 1, and the maximum is 329 ratings. The most rated movies (and not so coincidentally) the highest rated movies were:

- Forrest Gump (1994) - 329 ratings
- Pulp Fiction (1994) - 317 ratings
- The Shawshank Redemption (1994) - 307 ratings




Fig. 5: Distribution of the number of users and the amount of ratings. The first bin correspond to 20 movies rated.

IV. CLASSIFICATION MODELS

The three models here presented are first and foremost linear regression problems, where by fitting a model with the known data to the users and the movies, it will be possible to estimate the ratings of movies that have not been reviewed. In this sense, besides the model itself, it is necessary to fit two hyperparameters, the learning rate α , and the cost parameter λ . For both CF-LR and CF-LRC 20 features were used for users, and for movies (40 features in total). The number of features was defined empirically, based on the number of genres in the dataset (20 genres).

The first model developed aims to fit both users (x parameter) and movies (θ parameter) simultaneously, as depicted in the equation below, by optimization of the cost function with regularization.

$$\min_{x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{(i,j): r(i,j)=1} \left((\theta^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n \left(x_k^{(i)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n \left(\theta_k^{(j)} \right)^2 \quad (8)$$

The second model fits only the users parameter, given that it

accounts for the movies parameters by using the movie genres.

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} \left((\theta^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n \left(\theta_k^{(j)} \right)^2 \quad (9)$$

A. CF-LR

Different ranges of α (between 0.0001 and 0.002) and λ (between 0 and 100) were given to the model to be fit by usin 8 cross-validation in the training dataset. This allowed to progressively lower RMSE, throughout 500 iterations.

The best hyperparameters for Model01 were $\lambda = 6$ and $\alpha = 0.0005$, with an average RMSE from 8-CV of 1.25622.

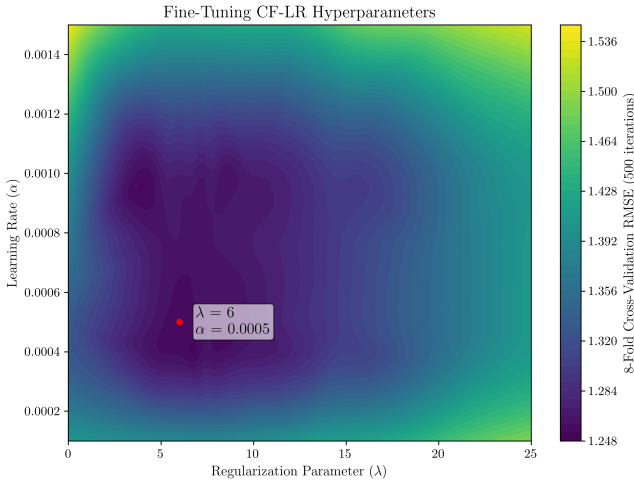


Fig. 6: Learning rate (α) and regularization parameter (λ) fitting.

Once the optimal hyperparameters were defined, the learning curve was estimated from the 8 CV on the training dataset, to assess the validity of the model and that it is not overfit. To further confirm that little to no overfitting is observed, in Fig. 7 the progression of the training score and the CV score tends to improve across the train set size, not overlapping or changing the trend significantly. It is important to bear in mind that there is a risk of slight overfitting, considering that there are 40 features to be fit.




Fig. 7: Learning curve, RMSE progression throughout the training dataset.

Using the defined hyperparameters, the cost function was estimated, where we see it sharply decreasing as the iterations run, finally converging at the limit.




Fig. 8: Cost function evolution with the number of iterations in the training set.

Once the model was finalized the train and test sets were compared for the errors (RMSE and MAE), as depicted in Table III, showing (as expected) the increase of errors in the test set in regards to the training set. The fact that errors increased significantly from the training to the testing set further supports the likelihood of overfitting the data.

TABLE III: CAPTIO CAPTION CAPTION

Dataset	RMSE	MAE	Support
Train Set	0.58842	0.44604	80668
Test Set	1.24037	0.90651	20168

B. CF-LRC

aqui é igual ao primeiro modelo mas ao inves de termos o x? (movies features) vamos ter os generos dos filmes, por

exemplo se um filme é romance/drama é [0,0,...,5,.5] tipo
 ns q mesmos dados de treino e de teste de smp
 neste caso temos o X q n precisa de otimizacao, pq é
 usado os generos dos filmes como features dos movies, content
 based, ent so vamos otimizar o theta (20 features dos users)
 entre alpha 0.007 e 0.0005 e lambda 10 e 0, foram otimizad-
 os os hyperparametros, como se pode ver na figura uma parte
 desse intervalo, para melhor visuazalicao da otimizacao, tendo
 sido esta com minimizacao do rmse e 8cv nos dados de treino




Fig. 9: CAPTIO CAPTION CAPTION

The best hyperparameters are: (0.05, 0.005) with 8cv
 RMSE: 0.98617

de movo a verificar a validade do modelo, temos a learning
 curve, a ver se n ha overfitting




Fig. 10: CAPTIO CAPTION CAPTION

e parece bastante bem, pq ambas as retas se acompanham
 para alem disso tbm e importante a verificacao da conver-
 gencia do modelo, atraves da sua funcao de custo




Fig. 11: CAPTIO CAPTION CAPTION

converge e bem

atendo as medidas de erro:

TABLE IV: CAPTIO CAPTION CAPTION

Dataset	RMSE	MAE	Support
Train Set	0.84859	0.65997	80668
Test Set	0.97254	0.74605	20168

medidas proximas, no signs of overfitting

C. CF-SVD

TABLE V: SVD model hyperparameters search space.

Hyperparameter	Possible Values
$n_factors$	{5, 15, 30, 40, 50}
lr_all	{0.005, 0.01, 0.05, 0.1, 1}
reg_all	{0.02, 0.1, 1, 5, 10}

Number of Singular Values or Components: $n_factors$

Learning rate: lr_all

Regularization Parameters: reg_all

Best Hyperparameters: {'n_factors': 50, 'lr_all': 0.01,
 'reg_all': 0.1}

lallla learning curve

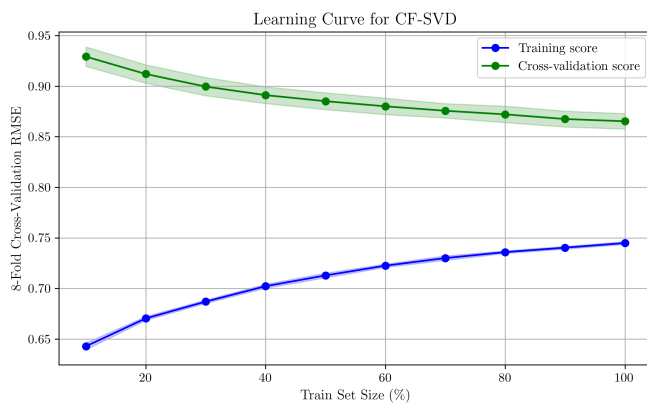


Fig. 12: CAPTIO CAPTION CAPTION

lalala results meauseres

TABLE VI: CAPTIO CAPTION CAPTION

Dataset	RMSE	MAE	Support
Train Set	0.75035	0.58313	80668
Test Set	0.87412	0.66957	20168

V. DISCUSSION

A. Performance Metrics




Fig. 13: caption caption caption

8cv aplicado ao dataset completo: 7folds para fit, 1 fold para teste (medidas do teste apresentadas), usando os modelos otimizados

dhuashdiuas

dhashdiaushda

dhauishdiaushd

ACHAS QUE VALE A PENA DEIXAR OS 4 GRAFICOS

???

PARA A OBTENCAo destes resultados

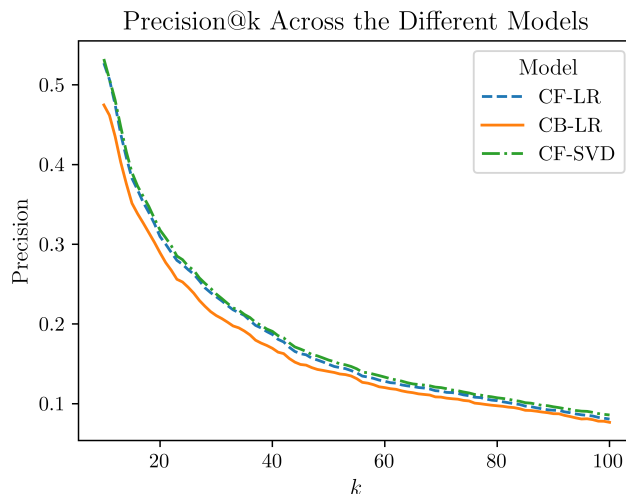


Fig. 14: caption caption caption




Fig. 15: caption caption caption

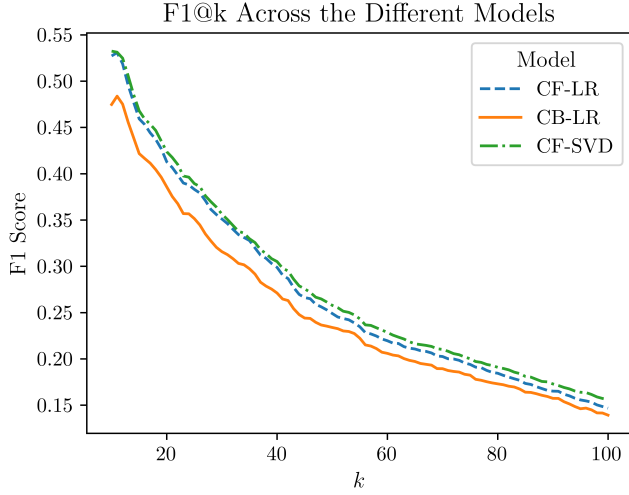


Fig. 16: caption caption caption




Fig. 17: caption caption caption

B. Literature Benchmark

TABLE VII: CAPTIO mesaures@100 across the models for the test data only

Measure	CF-LR	CB-LR	CF-SVD
Precision@100	0.08073	0.07659	0.08585
Recall@100	0.80732	0.76585	0.85854
F1@100	0.14679	0.13925	0.15610
MRR@100	0.46304	0.31212	0.38477
RMSE	1.23691	0.97240	0.87268
MAE	0.90417	0.74607	0.66804

VI. CONCLUSION

WORK LOAD

Both authors contributed equally to the project.

REFERENCES

- [1] J. A. Konstan and J. Riedl, "Recommender systems: From algorithms to user experience," *User Modeling and User-Adapted Interaction*, vol. 22, no. 1-2, pp. 101–123, 2012. [Online]. Available: <https://doi.org/10.1007/s11257-011-9112-x>
- [2] R. Katarya and O. P. Verma, "An effective collaborative movie recommender system with cuckoo search," *Egyptian Informatics Journal*, vol. 18, no. 2, pp. 105–112, 2017. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S1110866516300470>
- [3] A. Said and A. Bellogín, "Comparative recommender system evaluation: benchmarking recommendation frameworks," in *Proceedings of the 8th ACM Conference on Recommender Systems*, ser. RecSys '14. New York, NY, USA: Association for Computing Machinery, 2014, p. 129–136. [Online]. Available: <https://doi.org/10.1145/2645710.2645746>
- [4] F. M. Harper and J. A. Konstan, "The movielens datasets: History and context," *ACM Transactions on Interactive Intelligent Systems (TiiS)*, vol. 5, no. 4, pp. 19:1–19:19, 2015. [Online]. Available: <https://doi.org/10.1145/2827872>
- [5] Netflix Prize Community, "Netflix prize forum: Topic 1537," 2009, archived on the Wayback Machine. [Online]. Available: <https://web.archive.org/web/20090924184639/http://www.netflixprize.com/community/viewtopic.php?id=1537>
- [6] S. Funk, "Netflix update: Try this at home," 2006, accessed: 2025-01-09. [Online]. Available: <https://sifter.org/~simon/journal/20061211.html>
- [7] Y. Zhou, D. Wilkinson, R. Schreiber, and R. Pan, "Large-scale parallel collaborative filtering for the netflix prize," in *Algorithmic Aspects in Information and Management*, R. Fleischer and J. Xu, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2008, pp. 337–348.
- [8] P. Gopalan, J. M. Hofman, and D. M. Blei, "Scalable recommendation with poisson factorization," 2014. [Online]. Available: <https://arxiv.org/abs/1311.1704>
- [9] G. Research, "The movielens datasets," 2025, accessed: 2025-01-09. [Online]. Available: <https://grouplens.org/datasets/movielens/>
- [10] N. Hug, "Surprise: A python library for recommender systems," 2017, accessed: 2025-01-14. [Online]. Available: <https://surpriselib.com/>